

Response to Review: JCGS-25-420

Jayani P. Gamage Dianne Cook Paul Harrison
Michael Lydeamore Thiyanga S. Talagala

2025-10-25

We thank the anonymous associate editor and reviewers for their careful reviews. What follows is our point-by-point response to reviewer comments. The original reviewer comment is in normal font and *our response is in italic text*.

Note: RMSE has been changed to HBE throughout the paper. See Reviewer 2 comments for the reason.

Reviewer 1

Summary:

This manuscript proposes an algorithm to evaluate nonlinear dimension reduction (NLDR) techniques. The authors identify some characteristics and limitations of NLDR and develop an algorithm to analyze NLDR results' quality of how well they match the original high-dimensional data.

Strengths:

1. This manuscript is well-organized, providing sufficient background knowledge and examples which makes this paper easy to follow.
2. This paper contains sufficient technical details in the introduction of the algorithm.

Weakness:

1. The main concern is whether the proposed algorithm can properly evaluate the NLDR's performance. Using the RMSE plot with different binwidth values (e.g., in Figure 8 and Figure 11) to evaluate NLDR's performance is problematic for not considering the density within each cluster. For NLDR results that have denser clusters, for example layout a,c,g,h in Figure 11, a smaller binwidth needs be used to achieve the same average count in each bin, which will cause their lines in Figure 11 to be more upper left. However, denser clusters are not worse NLDR results.

Thank you for this thoughtful comment. This is a problem when comparing the HBE curves for different layouts because the same binwidth could be correspond to very different average bin density. This is something that we have discussed at length in developing the methods, and decided that using binwidth as the horizontal variable is more faithful comparison than using average bin density, primarily because it is directly controlled by the number of bins. We have revised the text in the paper to make this clear, and to note that the software allows for plotting against average bin density instead of binwidth, when comparing layouts from different methods and parameters. Other diagnostic measures, proportion of non-empty bins, relative average bin density and standardised bin count should also be considered in deciding the best fit, as illustrated.

2. In Figure 7, as RMSE steadily increases with the binwidth, it's not clear why the three binwidths 0.03, 0.05, and 0.07 are chosen. In addition, we shouldn't use a single simulated dataset to determine the default hyperparameters without any real analysis.

Thank you for raising this important point. The three binwidths (0.03, 0.05, and 0.07) represent small to moderate values, to illustrate exploring trade-offs between HBE, bin density and proportion of non-empty bins when decding which to use. This has been made clearer in the text.

The use of a single simulated dataset to illustrate a method is common practice. It needs to be sufficiently simple to communicate the method. The applications section has two examples that illustrate usage of the technique more generally. We have made this clearer in the text, when the simulated data is introduced.

3. For several figures that contains results from multiple NLDR algorithms, it would be nice to add the information of which algorithm/ what hyperparameter values for each of the subfigure. Many readers may have tried several NLDR algorithms themselves, and they will benefit from those figures to gain insights about behaviors of different NLDRs. Also it's nice to mention which data is used to generate each figure (e.g., in Figure 7).

Agreed. We have provided a table of the method and hyper-parameters used for every layout shown in the paper, in the appendix.

4. In section 3.6.1, it says “Values of b_1 between 2 and $b_1=\sqrt{n/r_2}$ are allowed”, why the choice of b_1 should consider r_2 ?

The aspect ratio of the NLDR layout is typically preferred to be conserved when plotting, and hence when we bin the data. r_2 reflects this aspect ratio, because r_1 is fixed at 1, for ease of calculations. The range of recommended values for b_1 is adapted from the Diaconis-Freedman Rule, where the upper limit takes aspect ratio into account. It is just a recommendation. Practically, 2 is too few bins, and the upper bound is possibly too many. We have revised the text to make this clearer.

Reviewer 2

Summary

The manuscript proposes a novel, mostly graphical diagnostic to evaluate nonlinear dimension reduction (NLDR) layouts by effectively “lifting” the 2-D embedding back into the original high-dimensional space and inspecting the fit visually via tours (i.e., sequences of 2D-linear projections of the original high-D space). The goal is to help analysts discern which 2-D projection is the most “honest” (faithful) representation of the true structure in the data. The paper’s contributions are:

1. an algorithm that overlays a geometric model on an NLDR layout using hexagonal binning and Delaunay triangulation in the embedding, then maps this model into the original large-D data space for visual assessment;
2. a new goodness-of-fit metric, the RMSE in original space, computed as the root mean squared distance between each point and its bin centroid in high-D (Eq 2) (3) interactive visualization tools (implemented in published software) to compare/investigate different NLDRs

It is mostly well written, addresses an important problem with critical real-world relevance, and provides a very well done, extensively documented R implementation. However, this reviewer is skeptical about the suitability of the newly proposed GOF metric to choose the most “reasonable” or “faithful” NLDR, finds the empirical evaluation of it to be severely lacking without comparison to alternative measures from the literature, and would request some restructuring and shortening of the paper.

We have now included a comparison with existing metrics, which shows that it performs favorably, and is a suitable addition to the analyst’s toolbox.

Strengths

Relevance, Accessibility, and Intuitiveness

The way this method visualizes the spatial distortions induced by NLDR seems uniquely accessible to me, and, coupled with the modern interactive/animated visualizations implemented for it, seems likely to become a very useful tool to investigate and understand high-D data structures while avoiding over-reliance on / misinterpretation of misleading 2D patterns so prevalent in many areas at the moment (e.g. [Izarry \(2024\)](#), [Chari & Pachter \(2023\)](#)).

Thank you. We have added these references.

Background Material & Examples

The paper is fully fleshed out with illuminating examples, and I commend the authors for the effort they put into documenting their package and providing very informative videos etc of the method in action.

Reproducibility and Code Availability

The paper is exemplary in this regard - the underlying software is a CRAN compliant R package, and the Github repository itself allows to recompile the entire paper and appendix with all figures etc from scratch.

Major Issues

No Standard Embedding Quality Metrics or Comparisons

My primary concern is the absence of established quantitative metrics

1. to evaluate to what extent the newly proposed “RMSE” metric corresponds to established notion of embedding faithfulness and
2. to validate the different embeddings under consideration.

The paper’s only evaluation of “honesty” is done via the proposed custom “RMSE” in high dimensional space (defined in Section 3.4) and visual heuristics, but it does not report any conventional measures like trustworthiness/continuity scores, k -Neighborhood preservation rates (as derived from the co-ranking matrix - e.g. R_{NX} -curves for different neighborhood sizes k , c.f. ([Lee et al, 2015](#)), ([Lueks et al, 2011](#))), or Shepard diagrams of original space vs. embedding space distances. This is a significant omission because NLDR almost always faces a trade-off between optimizing preservation of local structures vs preservation of global

structures, which these metrics are (mostly) designed to resolve / analyze. For example, the authors could have computed such measures of trustworthiness for each NLDR layout to see if the one with lowest “RMSE” indeed had the highest (local and or global) neighborhood structure preservation but no such analysis was provided. The manuscript doesn’t even mention or discuss these metrics, which is not acceptable in a study focused on faithful representation, in my opinion.

Questions a revised version should aim to resolve:

1. does “RMSE” measure local or global structure preservation (or a mixture of both)? From the definition, it seems to be focused on how well small neighborhoods in embedding space (the bins) correspond to small neighborhoods in high-D space, with, IIUC, particularly high values if there are lots of “hard intrusions” (i.e., points that are embedded in the same 2D bin despite being far apart in high-D space)?

We thank the reviewer for this insightful question. The HBE metric is intended to complement existing measures rather than replace them, and its sensitivity depends on both the bin resolution and the spatial distribution of the NLDR layout. When fine-grained bins are used, HBE primarily captures the preservation of local neighborhoods, reflecting how well small groups of points in the NLDR layout correspond to their high-dimensional counterparts. For coarser bins, the metric incorporates more global structure, effectively averaging over larger regions of the NLDR layout.

2. how strongly are “RMSE” values correlated to other established measures of (local or global) structure preservation? and (how) are these correlations affected by the binning hyperparameters?
3. are the NLDRs selected by lowest “RMSE” the ones that would also be selected by other, more established measures of embedding faithfulness whose characteristics are well understood? at the very least, such measures should be provided, compared and discussed for all the embeddings under consideration in the case studies.

We thank the reviewer for this insightful question. HBE is intended as a diagnostic tool rather than a definitive selection metric. Its primary purpose is to provide an interpretable measure of how well high-dimensional structure is captured in the NLDR layout. Traditional measures such as $AUC(R_NX)$, RTA , RGS , and sc evaluate local or global structure preservation, but their values can vary substantially depending on method settings. In contrast, HBE offers a more connection between the high-dimensional data space and the fitted model in the embedding space, highlighting how faithfully an NLDR method captures the underlying data structure.

In the revised manuscript (Section 4), we include a parallel coordinate plot for the 2NC7 data comparing HBE and existing metrics. The appendix contains parallel coordinate plots for the PBMC3k and MNIST digit 1 datasets. These figures show that HBE offers a different way of looking at how well the NLDR layouts works, while the other metrics capture related but with different settings.

The reliance of “RMSE” on euclidean distances in high-D space also seems limiting:

1. at least for coarser bins and/or highly non-linear data structures, geodesic distances along the data manifold à la ISOMAP are likely to measure a more relevant kind of dissimilarity
2. a definition based on a more general notion of distance would make this applicable to investigating 2D embeddings of data types for which the euclidean distance of their numerical data representations is not a suitable/applicable measure of dissimilarity (e.g. images, text).

We thank the reviewer for raising this important point. While Euclidean distance may not be ideal for highly nonlinear manifolds or complex data shapes, it can still provide a meaningful summary when the high-dimensional means used in HBE are suitably positioned. In such cases, Euclidean distances effectively capture variation from the fitted model, allowing HBE to quantify how well the NLDR layout preserves the structure of the high-dimensional data.

Insufficient Validation of the Proposed Diagnostic (RMSE) Across Methods

Related to the above points - the paper has no theoretical or empirical demonstration that minimizing this “RMSE” aligns with minimizing other distortion measures like stress or maximizing neighborhood preservation (as measured by the co-ranking matrix). Furthermore, the claim that RMSE can be used to compare different NLDR methods on equal footing deserves scrutiny. Different methods optimize different criteria; it’s plausible that one method might (tends to) achieve lower “RMSE” simply by construction, especially if my reading that it mostly rewards embeddings that manage to avoid hard intrusions is correct. Overall, much more critical validation of the RMSE diagnostic is needed to convince readers that “honesty” as defined here aligns with meaningful goodness-of-fit in embedding.

We appreciate this thoughtful comment highlighting the need for more rigorous validation of the HBE diagnostic across NLDR methods. We agree that different NLDR algorithms optimize different objectives and that a single quantitative diagnostic should not be interpreted as a universal criterion of embedding quality. Our intent is not to claim that HBE replaces established distortion or neighborhood preservation measures, but rather to propose it as a complementary diagnostic that provides a direct and interpretable link between the high-dimensional model and its low-dimensional representation.

To clarify this point, we have revised the text to explicitly state that HBE is not intended to be minimized across all methods in a comparative sense, but instead used as a diagnostic to evaluate the faithfulness of the mapping between the fitted model in the high-dimensional space and the observed embedding. We also acknowledge that empirical validation against stress and co-ranking based measures would be a valuable avenue for future work, and we have added a statement in the discussion noting this as a natural next step.

Limitations of the Evaluation (Simulation Study Design):

The manuscript includes a simulation to illustrate the approach (the “2NC7” dataset with two nonlinear clusters in 7-D) and two real-data applications (single-cell RNA-seq and an image dataset of handwritten digits). While these examples are appropriate, the scope of simulation scenarios is narrow, and the analysis lacks any notion of variability. The simulation seems to be a single realization of a specific two-cluster scenario. The authors did not explore, for instance, varying the overlap between clusters, adding more clusters or different manifold shapes, or introducing different noise levels. Ideally, the authors would systematically vary factors like: number of clusters or pattern types, intrinsic dimensionality of structures, sample size, etc., and show that their approach reliably selects the truth (or at least, the most interpretable view) in each case, and how well alternative measures would do so. Because only one run is shown, there is also no information how stable the RMSE curves in Figures 8 and 13 are across replications. I encourage the authors to perform more systematic, replicated simulations, report measures of variability for their performance evaluations, compare against relevant alternatives, and take much more care not to overstate their conclusions beyond the tested scenarios (e.g. the claim that it “identify bad representations so they can be avoided” is based on a very limited set of examples).

We thank the reviewer for this constructive feedback and agree that the current simulation study represents a limited scenario. A broader evaluation varying factors such as cluster separation, noise level, number of clusters, and intrinsic dimensionality would provide stronger validation of the diagnostic. Our goal in this paper, however, was to introduce and illustrate the framework rather than conduct a full-scale benchmarking study. The 2NC7 dataset was designed to demonstrate the diagnostic’s behavior in a controlled nonlinear setting with a known ground truth, enabling clear interpretation of how it identifies misrepresentations.

Minor Issues

1. First part of the title seems inappropriate for a scientific paper to me, as does the implied (and not substantiated, IMO) claim that the proposal is (primarily) useful to reliably select one of many NLDRs (as opposed to primarily useful for visually investigating details about a specific 2D NLDR).

We thank the reviewer for this insightful comment. We agree that the original title was informal and could be misinterpreted as over-stating the scope of our work. We have revised the title to focus on the main contribution. The revised title is: “Beyond Pretty Pictures: New Visual Tool for Investigating Nonlinear Dimension Reduction”.

2. For marketing and clarity reasons, I would strongly suggest to not name the metric “RMSE”, but something more distinctive/descriptive.

We agree that the name “HBE” may cause confusion with the standard Root Mean Square Error metric. To improve clarity and better reflect the purpose of the measure, we have renamed it to HBE (hexbin error) throughout the manuscript. This new name more accurately conveys that the metric quantifies structure preservation rather than prediction error.

3. Flow from 3.1 to 3.2 (with too many sub-subsections) is difficult. Sections 3.2.1 (Scale the data) and 3.2.2 (Construct hexagon grid) are just short paragraphs, would be clearer if combined into a single numbered list of steps for the model construction. Much of 3 might be clearer and more compact if simply replaced by a direct step-by-step description of the algorithm in algorithmic pseudocode form with suitable short explanations.

We thank the reviewer for this helpful suggestion. We agree that summarizing the algorithmic flow is important for readability. However, we have chosen to retain the current structure with distinct subsections to maintain conceptual clarity. Each subsection (e.g., scaling, grid configuration, binning, and neighborhood construction) describes a logically separate stage of the model-fitting process that contributes to the interpretability of the proposed framework. This modular structure also helps readers link the algorithm steps with their mathematical foundations and corresponding figures.

4. Section 3.5 is a claim without supporting evidence – many out-of-sample embedding methods exist (some based on rather similar ideas, which should be referenced), you would need to show that this particular one actually works reasonably well for this to deserve space in the paper. I would strongly suggest to simply mention this in the discussion section as a possible avenue for further development instead.

We thank the reviewer for this comment. We agree that the predictive mapping described in Section 3.5 would require further empirical evaluation to substantiate its performance and to clarify its relation to existing out-of-sample embedding methods. We have therefore revised Section 3.5 to clarify that this feature is a conceptual extension rather than a tested result, and we have added citations to related work such as PCA, autoencoders, UMAP’s transform function, and parametric tSNE. This revision places the idea in a broader methodological context and highlights it as an avenue for future research rather than a core contribution.

5. Section 3.6 should be drastically shortened and details moved into an appendix.

We appreciate the reviewer’s feedback. We agree that keeping the section clear and concise is important, and we have carefully reviewed Section 3.6 with this in mind. However, we have retained this section in the main text because the tuning procedure is essential for model interpretability and reproducibility. The parameters described (e.g., number of bins, bin density cutoff) directly affect the model fit and are fundamental to achieving an accurate representation. Keeping this section ensures that readers can understand how the tuning process influences the process of model fitting.

6. p. 18: “A particular pattern that we commonly see is that analysts tend to pick layouts with clusters that have big separations between them.” The literature/internet contains many lively discussions of the perils of mis-interpreting and post-hoc’ing NLDRs (e.g. (Izarry (2024)), (Chari & Pachter (2023))), it might be good to cite/summarize these instead of/in addition to basing your motivation on personal anecdotal evidence? The tone here seems overly conversational to me (also: “we almost always see there are no big separations in the data”) and should be made more formal.

We thank the reviewer for this insightful comment and agree that this paragraph benefited from a more formal tone and stronger connection to the literature. We have revised the section to reference recent discussions highlighting the risks of misinterpreting NLDR results (Izarry, 2024; Chari & Pachter, 2023). The revised paragraph now adopts a more objective tone and situates our argument within the broader context of current methodological discourse. Our intent is to emphasize the importance of evaluating NLDR results critically, using evidence-based references rather than anecdotal experience.

7. Section 4, esp. p.19: IMO, “to compare and assess a range of representations” an analyst really should look at their respective RNX curves, stress values, and Shepard diagrams, think about whether they care more about local or global structure preservation, and then decide based on those factors. The method proposed here seems to me to be mostly useful to then investigate specific regions / slices of high-D space to see for which sets of data points the high-D structure is (not) preserved well in the NLDR. I find this whole section to be vastly overselling the proposal – it ignores and implicitly discards most of the prior work in this field. Its intro paragraphs are also highly redundant with other content in the paper and should be cut.

We appreciate the reviewer’s thoughtful comment and agree that RNX curves, Shepard diagrams, and stress values are valuable diagnostics for assessing NLDR performance. Our intent, however, is not to replace or disregard these established measures but to provide a complementary diagnostic framework that directly links the low-dimensional representation to the underlying high-dimensional structure through a model-based approach. Traditional metrics primarily evaluate structural preservation (local or global) within the embedding, whereas our HBE-based approach quantifies how well an interpretable model in p -D corresponds to the observed 2-D layout.

This distinction is important because analysts often rely on NLDR visualizations to make structural interpretations (e.g., clusters, manifolds), yet the existing metrics may not indicate where those interpretations fail locally. Our framework extends the diagnostic toolbox by offering both global (HBE) and local (bin-level residual) assessments that can be used alongside RNX, Shepard, or stress-based analyses. We have revised the text to clarify that our method is intended as a complementary approach for identifying and interpreting poor or misleading embeddings rather than as a replacement for existing evaluation techniques.

8. Section 5 should be moved into an appendix to shorten the paper somewhat. Figure in Section 5.1. could benefit from coloring the clusters in different colors so global structure preservation is also somewhat legible from the figures for the NLDRs (e.g. “is orange cluster between green and blue as in high-D or do they switch positions”). Section 5.2 might also have benefited from a corresponding ISOMAP or similar embedding of this data suitable for truthfully (isometrically!) unrolling the high-D manifold, and then showcasing that such a successful “unrolling” can be discerned from the way the projections look? tSNE just isn’t isometric, so the distortions visible here are entirely expected, IMO?

We thank the reviewer for these helpful suggestions. We have carefully considered moving Section 5 to the appendix; however, we believe it is important to keep it in the main text because it demonstrates the interpretive value of our proposed method. The examples in Section 5 illustrate how examining the model in the data space reveals structural distortions introduced by NLDR methods an essential contribution of our approach that directly supports the paper’s central argument. Moving it to the appendix would significantly reduce the clarity of this methodological insight for readers.

To address the reviewer’s second suggestion, we have updated Figure 5.1 to color the clusters distinctly. This enhancement makes the global relationships between clusters more visible and allows the reader to better assess the preservation (or loss) of structure across NLDR methods.

We agree that an isometric method such as ISOMAP could, in theory, produce a more faithful “unrolling” of the high-dimensional manifold. However, the purpose of Section 5.2 was not to determine which NLDR method performs best, but rather to demonstrate how uneven data density can distort the resulting low-dimensional layouts, even under commonly used non-isometric methods like tSNE. We have clarified this motivation in the text and added a note acknowledging that ISOMAP could serve as a complementary example of a more isometric representation. Our focus, however, remains on diagnosing distortions and understanding their implications for interpretation rather than comparing algorithms.

Typos and grammar:

Revised version should be much more carefully proofread: “appropraite”, “doen’t”, “it is has two separated nonlinear clusters”, “the methods tNSE”

We appreciate your careful attention to detail and for pointing out the typographical and grammatical errors. We have carefully proofread the entire manuscript and corrected the issues you identified (e.g., “appropraite” → “appropriate,” “doen’t” → “doesn’t,” “hastwo” → “has two,” “tNSE” → “tSNE”), along with several other minor typographical inconsistencies. The revised version has been thoroughly checked to ensure clarity and correctness.